ikat: weaving Markdown into LaTeX

opticsWolf

October 26, 2026

Abstract

Manuscripts live in Markdown — camera-ready papers live in LaTeX. ikat closes that gap with a strict-subset Markdown pipeline: a Rust core owns parsing, layout math, and code generation while thin Python wrappers carry strings across the boundary. Float placement, publisher templates, and engine choice are configuration, not post-processing. This paper is its own demo: every heading, float, table, citation, and plot below was woven by ikat itself. (118 Rust tests, 76 Python tests, and compile proofs guard the behavior; acceptance is one document rebuilt to the byte.)

1 Introduction

The journey from draft to submission is a format translation with opinions. Authors think in sections, figures, and `code`; publishers think in floats, baselines, and bibliographies. Between them sits a manual step of paste, reformat, and renumber that is tedious, brittle, and repeated for every revision. Tools that automate it tend to fail in one of two directions: they accept all of Markdown and produce fragile LaTeX, or they emit beautiful LaTeX from a language nobody writes. ikat takes the third path: a strict subset of Markdown with **fail-fast** errors, deterministic output, and a golden master test that keeps the published paper byte-identical across refactors.

The architecture follows one rule: Rust owns compute, Python carries strings. Parsing, float resolution, TikZ emission, and pgfplots generation live in the `ikat` crate; the `ikat` package exposes `build_document`, a template library, a CLI, and an MCP server over a single `BuildSpec` JSON boundary. Nothing crosses that boundary except data (no callbacks, no shared state, no surprises). The whole pipeline is $O(n)$ in document size: events in, LaTeX out, with no tree that could blow up on a long manuscript.

Three doctrines shape every decision. First, the subset is strict: unknown attributes, dangling citation keys, and doubled skeleton tokens are build errors. LaTeX logs are never the debugger. Second, defaults never move: tuning emits preamble lines only when they differ from the LaTeX defaults, so untouched documents weave byte-identically. Third, the paper is the test suite, and §4 names the numbers. Sections that follow use mechanisms §2–§5 describe; §6 collects the skeleton token contract.

The paper you are reading is the demo and the test at once. Its four diagrams, four tables, and two plots were woven from the Markdown source beside this PDF: no hand-placed floats, no touched-up `.tex`. Two of the diagrams sit anchored in their paragraphs (`span=column`

`pos=here`, scaled to the column with `width`); Figures 3 and 4 share a dedicated float page (`pos=page`). §3 walks that choice as a worked example. Nothing here was pasted.

2 Architecture

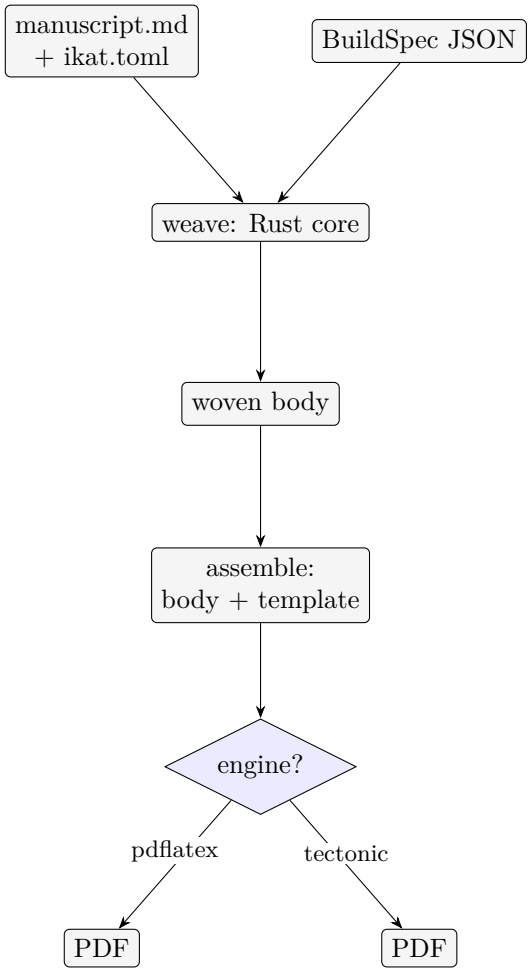


Figure 1: The ikat build: sources in, assembly, one engine decision.

Figure 1 shows the build: sources in, assembly, one engine decision. The manuscript carries content plus per-element {`attrs`}; `ikat.toml` carries document policy: which kinds span columns, where floats may go, which template wraps the result. The spec carries per-build registries: diagram captions, plot captions, table captions, citation keys. The core weaves the body, the template wraps it, the engine compiles it. Each stage validates its own inputs: unknown {`pos=center`} dies in Rust with the fence number, a template that drops a needed package dies naming it, a missing binary dies with the install command. Nothing fails in a log. Figure 1 is anchored where it is declared, so it reads as part of this section rather than floating to the page top; Figures 3 and 4 share the float page after this section (`pos=page` queues both for the same dedicated page): tall floats that queue together land together.

The subset boundary is deliberate. The block splitter runs on pulldown-cmark the CommonMark pull parser [1], which owns block boundaries while hand-written rules still parse every fence, table, and directive from raw source slices; the published paper rebuilds byte-identically under either splitter. Our syntax (bracketed author-key citations, {`span=`} attrs, `%%` directives) stays custom, and constructs outside it feed their raw lines back as paragraph text instead of failing. Mermaid itself is a strict subset too: flowcharts with `-->`, `---`, edge `|labels|`, and `[]`/`{}` nodes. Anything else is a string error naming the line, because a diagram that almost renders is worse than one that refuses.

Figure 2 replays the same build as a sequence: the calls §2 names (`convert()` across the PyO3 boundary, `compile_pdf` driving the engine), with the engine choice as an `alt` box instead of prose. Like Figure 1 it is anchored in its paragraph and scaled to the column.

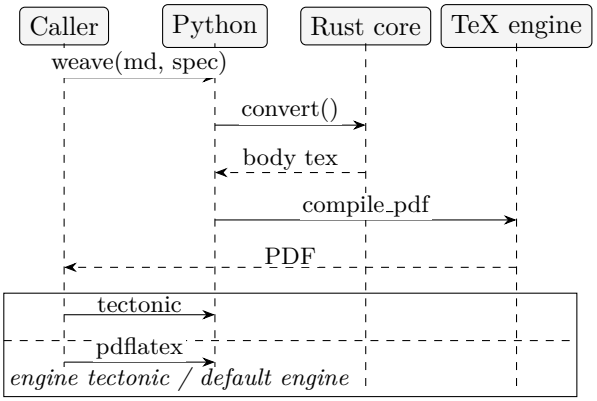


Figure 2: The same build as a sequence: calls across the PyO3 boundary, engine choice as an alt box.

Sequence diagrams dispatch on their header line through the same `flowchart_to_tikz` entry point: participants declare columns in order, `->>` and `-->>` messages take one row each, and `alt`/`else`/`opt`/`end` boxes span all columns. The subset is strict like the flowchart one: `loop`, `par`, `Note`, and self-messages are string errors naming the statement.

Figure 3 shows the document path as states: fences convert, the body wraps and validates, the engine compiles, with the weave stages grouped in one composite box. State diagrams take the same entry point (`stateDiagram-v2` header): `[*]` markers, `-->` transitions with optional labels, `state "Label" as Name`, and one level of `state Name { ... }` boxes drawn around laid-out members, so a box can never move coordinates. Nested composites are a build error, like everything else outside the subset.

Figure 4 maps the same policy as a flowchart: element kind and {`attrs`} resolve against `[spans]` and `[floats]`, and each combination loads only the packages it needs. It shares a

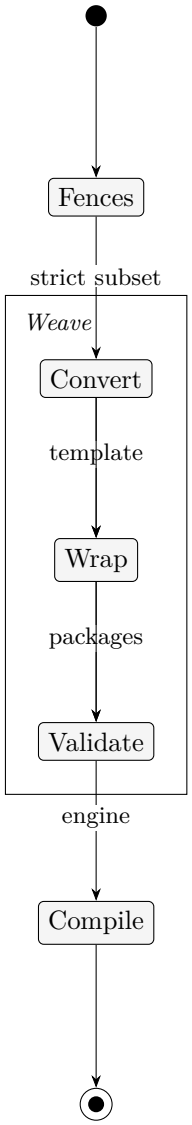


Figure 3: The document path as states: fences convert, the body wraps and validates, the engine compiles.

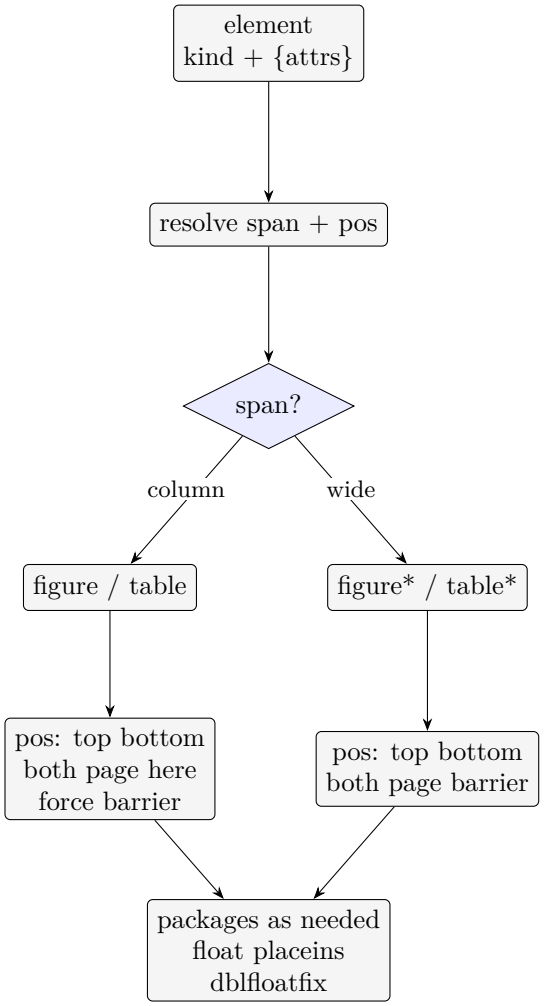


Figure 4: Float resolution: attrs against policy.

float page with Figure 3 because both queue with `pos=page`.

3 Floats without fear

LaTeX float placement is where manuscripts go to drift: a figure declared `[t]` lands three pages later and nobody knows why. ikat treats every floatable element (diagram, picture, plot, table) uniformly: per-element attributes resolve against document policy, and the LaTeX rules become build errors instead of log mysteries. A `figure*` with `here` or `force` is illegal LaTeX, so it is illegal here too, with the fix (`span=column`) in the message. Wide floats at the bottom load `dblfloatfix` automatically; `force` and `barrier` load `float` and `placeins`, or require them in template heads. Table 1 lists the attribute surface.

Table 1: Per-element float attributes.

Attribute	Values	Default
span	column or wide	per-kind `[spans]`
pos	top, bottom, both, page, here, force, barrier	`[floats]` default
width	fraction of span or TeX length	span width
captionpos	top or bottom	figures bottom, tables top

Figures 1 and 2 above demonstrate the anchored end of the range: column figures placed `here` and scaled with `width`, declared next to the paragraphs that reference them.

Document-wide tuning lives in `[floats]`: this paper sets `topfraction` 0.9, `bottomfraction`

0.7, `textfraction` 0.1, `topnumber` 3, `bottomnumber` 2, and `pos_default` top: six lines that are the only float preamble it emits. The numbers are LaTeX counters, not suggestions: `topnumber` caps how many floats may share a column top, `textfraction` guarantees text at least a tenth of every text page, and a float that fits nowhere waits in the queue instead of squeezing in. Short floats referenced mid-paragraph (like Tables 1 and 2) take `pos=force` ([H]) and never queue at all; tall ones (Figures 3 and 4) take `pos=page` and share one dedicated float page. Everything else is LaTeX default, so the tuning stays visible and minimal. `barrier_sections` is on in this paper: each section starts with its predecessor's floats flushed, so no heading strands pages from its text. The engines table in §5 adds the per-element `barrier` on top, pinning itself above its own section.

4 Evidence

Suite	What it guards	Count
Rust unit	emitters, scanner, templates, floats	118 run
Python	pipeline, CLI, MCP, errors, golden paper	76 pass
Compile proofs	heads, skeleton, floats, engines	PDF-verified

Table 2: Test suites and what they guard.

The suites in Table 2 overlap on purpose: Rust tests pin units, Python tests pin the surface, and compile proofs pin reality, because a `.tex` that never met `pdflatex` is a rumor. Every rejection names its cause: unknown statements echo the manuscript line and the offending text with a caret under the token, so new syntax fails loudly instead of drifting. The golden master guards the output side (the published paper rebuilds byte-identically); the error catalog guards the input.

Note: This paper is proof zero: woven, compiled, and text-verified in one command.

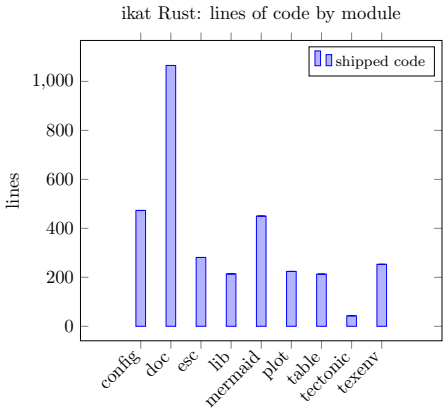


Figure 5: Rust lines of code by shipped module (wc, HEAD; test-only spike excluded).

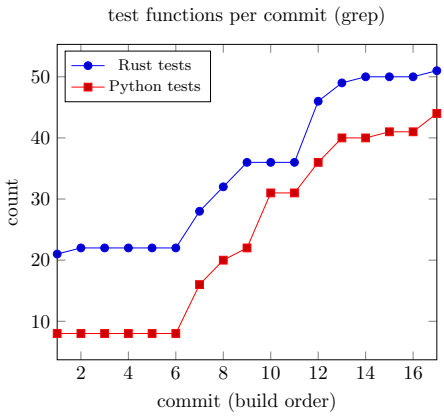


Figure 6: Test functions per commit, both suites (grep, no smoothing).

4.1 Where the code and the tests live

Figure 5 counts lines per shipped Rust module (`doc` dominates because assembly lives there, the test-only spike module aside), and Figure 6 counts test functions per commit across the build, Rust and Python series separately. Both series are grep-true: `#[test]` attributes and `def test` functions, counted from history, with no smoothing and no invention. The Python steps track texenv tests (7), the CLI/MCP surface (10), floats (12), and skeletons (13); the Rust steps track texenv (7), templates (8), heads (9), floats (12), skeletons (13), and the AST spike (14). Growth is linear because each milestone ships its proofs. The spike paper trail is typical: conjecture, probe, verdict, all in the tree.

5 Engines and templates

Table 4: Compile engines compared.

Question	**pdflatex (default)**	**tectonic (optional)**
Base engine	pdfTeX via TinyTeX	embedded XeTeX 0.17.0
Install	TeX Live tree + tlmgr	zero (binary or cargo feature)
arXiv fidelity	closest (same engine family)	snapshot drift possible
BibTeX	runs when a `.bib` is staged	auto-runs always
Needs from ikat	nothing special	drop `inputenc`, never emit empty bibliography

The default engine is pdfTeX because submissions are the authority: arXiv runs TeX Live, so the closest local engine wins and the rest is convenience. Tectonic the self-contained TeX engine [2] is the zero-install path (one binary, cached bundles, no TeX tree), and it caught two real bugs on its first run: XeTeX dies on `[utf8]{inputenc}`, which ikat now drops automatically, and an empty `\{}bibliography{}` is fatal under an engine that auto-runs BibTeX, so the core no longer emits one. Templates come in three levels: the generated head (level 0), nine shipped publisher heads (level 1 covers arXiv, article, IEEE the document class [3], ACM, LNCS, Elsevier, and APS), and whole-document skeletons (level 3, the escape hatch this paper uses). Level 1 heads carry {{{`title`}}} tokens; level 3 carries the contract in §6. Plotting needs no matplotlib: bar and line emitters write pgfplots directly the TeX plotting package [4], compiled standalone to the PDFs below. Engine choice is

Table 3: Skeleton token contract.

Token	Fills from	Required
{{title}} {{author}} {{thanks}}	manuscript H1, spec author and thanks	as level 1, braces included
{{body}}	woven floats, tables, and text	exactly once, always
{{bibliography}}	style plus staged .bib name	exactly once unless bib_name is empty
{{abstract}}	the whole woven abstract env	exactly once when the manuscript has one

per build, not per document: the same woven body compiles under either engine from one flag (--engine), and both PDFs ship the same page count and geometry (the showcase rebuilds clean under tectonic in CI). Templates compose with engines freely: no head names an engine, and no engine names a head. The cost of the second engine is one probe document and a checksum-pinned binary, both in the tree.

6 Appendix: skeleton token contract

The contract is small on purpose: four token families cover title, body, bibliography, and abstract, each filled exactly once (Table 3 above lists all four). {{thanks}} rides with {{title}} because only some heads set it; empty thanks emit nothing rather than an empty footnote. {{bibliography}} is conditional on bib_name: an engine that auto-runs BibTeX dies on an empty {{bibliography}} line, so the core omits it instead of emitting it. Unknown and doubled tokens are build errors naming the file and line; TeX comment lines never count as content. Skeletons are the escape hatch, not the default: nine publisher heads cover the common cases, and level 3 exists for the tenth. Submissions want sources, not PDFs: arXiv takes the .tex, the .bib, and every figure file, and rebuilds. The build directory holds exactly that set after one command: woven .tex, staged .bib, and standalone plot PDFs beside it. Nothing in the manuscript names an engine or a template, so the same sources submit anywhere a head exists.

References

[1] "pulldown-cmark: a pull parser for CommonMark in Rust," github.com/pulldown-cmark/pulldown-cmark, 2026, accessed 2026-10-07.

[2] "Tectonic: a modernized, complete, self-contained TeX/LaTeX engine," tectonic-typesetting.github.io, 2026, accessed 2026-10-07.

[3] "IEEEtran: IEEE LaTeX document class," ctan.org/pkg/ieeetran, 2026, accessed 2026-10-07.

[4] "pgfplots: a TeX package to draw plots," ctan.org/pkg/pgfplots, 2026, accessed 2026-10-07.